# Definitions
# What is ML ?

- ***a* definition:**

  For a given **Task** T, a **machine** (algorithm) A
  obtains better **performance** P
  after an **experiment** E. (It has ***learned*** from it)
  (Experiment ~ data)

**Yann LeCun's cake metaphor:**

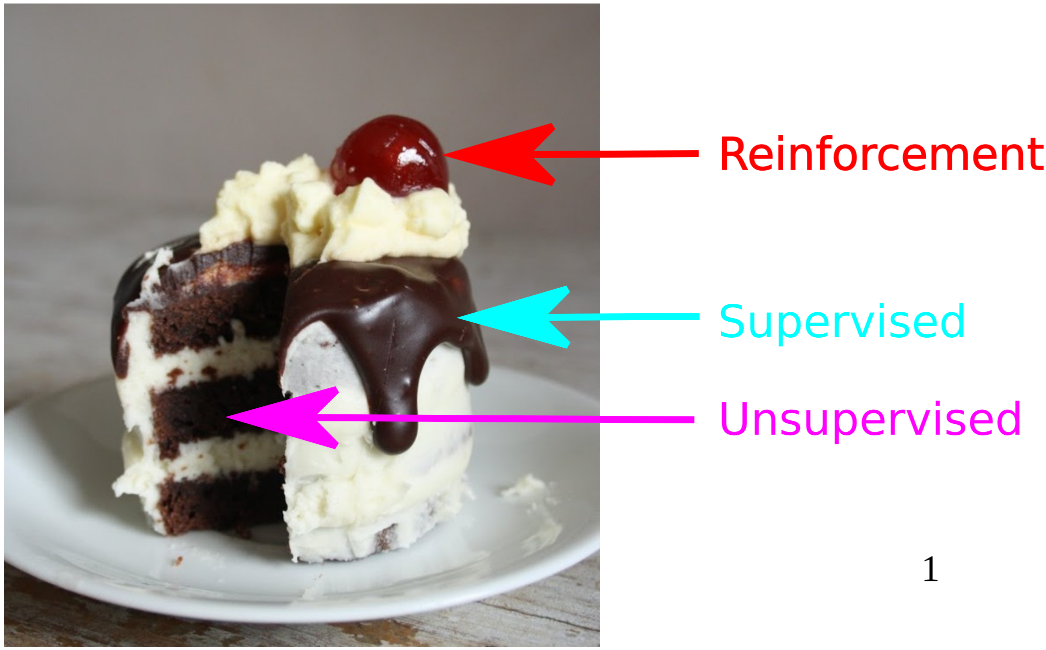


- **<u>3 types</u>** of learning :

  **- Supervised:** w/ labels

  **- Unsupervised:** w/o labels
  (incl. self-supervised)

  **- Reinforcement**
  (outside this course)

# Today – Outline

- **Supervised** Learning basics:
  - Linear **regression**
- Lots of **Vocabulary**, notations
- Optimization basics: **Gradient Descent**
- **Supervised** Learning
  - Classification with the Perceptron (maybe)

# Today:
# **Supervised Learning**

Input: $\vec{x}_n = (x_{n,d})_{d\in[1,\ldots,D]}, X = \{\vec{x}_n\}_{n\in[1,\ldots,N]}$

- Expected Output: $y^{GT}$ or $t_n$ (***Ground Truth***)

  Which kind of Task $\rightarrow$ depends on type of $t_n$

- ***Model***: $y^{predicted} \equiv \hat{y}_n = \sigma(f_\Theta(\vec{x}_n))$

  fct. $f_\Theta$ is ***parameterized*** by ***parameters*** $\Theta$

- ***Learning :*** finding *optimal* parameters to minimize discrepancy between $\hat{y}$ and Ground Truth $t$

$$\Theta^* = argmin_\Theta\Big(\sum_n^N \ell(\hat{y}_n, t_n)\Big)$$

- ***Cost** Function (**loss** function)* : to be chosen

# **Supervised** Learning: **Regression**

Pairs of data points $\vec{x}_n = (x_{n,1}, x_{n,2})$

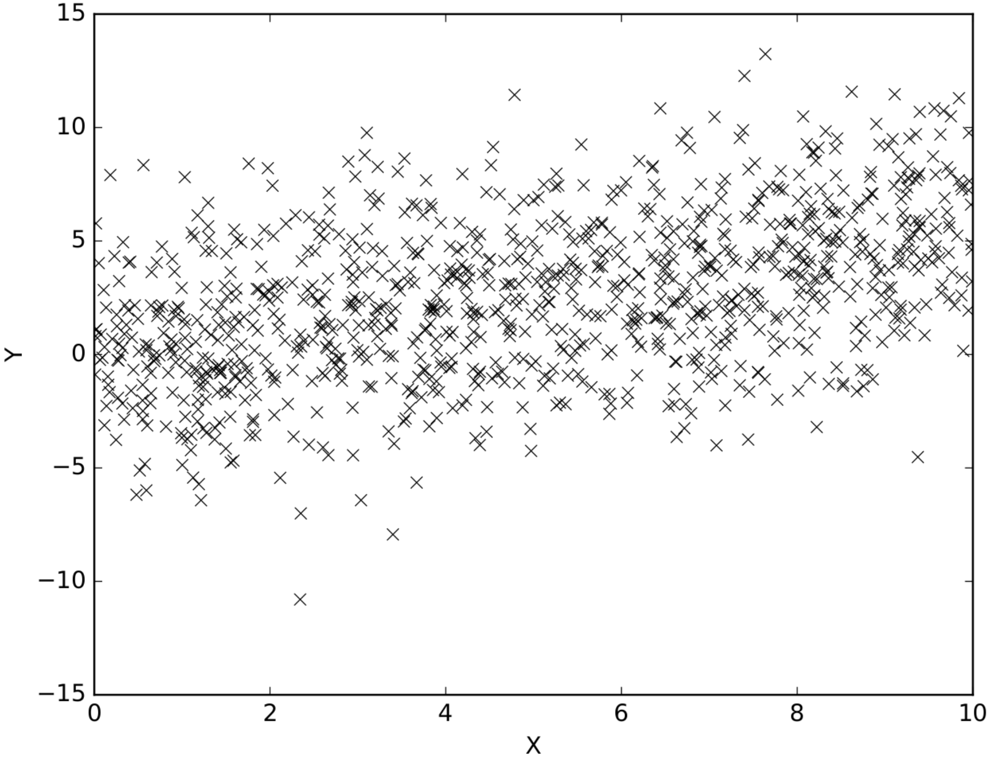


→ Relationship $f(x)=y$ ?

→ **Regression**

- ***linear***: $f_{a,b}(x) = ax + b$ or $f_{\vec{a},b}(\vec{x}) = \vec{a} \cdot \vec{x} + b$

- ***polynomial***: $f_{\Theta}(\vec{x}) = \vec{\theta} \cdot \Phi(\vec{x})$

*(degree $P$)* *(see later, feature maps)*

# More Vocabulary
## (+case of Regression)

Input: $\vec{x}_n = (x_{n,d})_{d \in [1,\ldots,D]}, X = \{\vec{x}_n\}_{n \in [1,\ldots,N]} = (x_{n,d})_{(N,D)}$

- *Ground Truth*: $t_n \in \mathbb{R}, T = \{t_n\}_{n \in [1,\ldots,N]}$

  **Continuous** output → Task is **Regression**

- **Model: e.g.** a linear function of the input : $f_{\vec{a},b}(\vec{x}) = \vec{a} \cdot \vec{x} + b$

  **-** Parameters: $\Theta = \{b, a_d; d = 1, .., D\}$

  - **Prediction:** simply $\hat{y}_n = f_\Theta(\vec{x}_n)$

  $Card(\Theta) = 1 + D$

- Learning **Algorithm:**

  - ***Initialization:*** $\Theta = \Theta_0$

  - Minimize some Loss $\ell(\hat{y}_n, t_n)$ (to be chosen)

  - For this, use some minimization scheme (Grad. Desc.)

# Supervised Learning: **Regression**

- We can choose: **Least Squares**

$$\text{Single data point Loss: } \ell(f_{\Theta}(\vec{x}_n), t_n) = (f(\vec{x}_n) - t_n)^2$$

$$\text{Gloabal Loss: } \mathcal{L}(\Theta, X, T) = \frac{1}{N}\sum_{n=1}^{N} \ell(f_{\Theta}(\vec{x}_n), t_n)$$

- **Gradient Descent**:

$$\vec{\theta} \to \vec{\theta} - \eta \vec{\nabla}_{\vec{\theta}} \mathcal{L}(\theta, X, T)$$

# Gradient Descent
## short reminder

- I have a function $J(\theta)$ and want to find the value $\theta*$ for which $J(\theta)$ is minimum

# What is the gradient ?

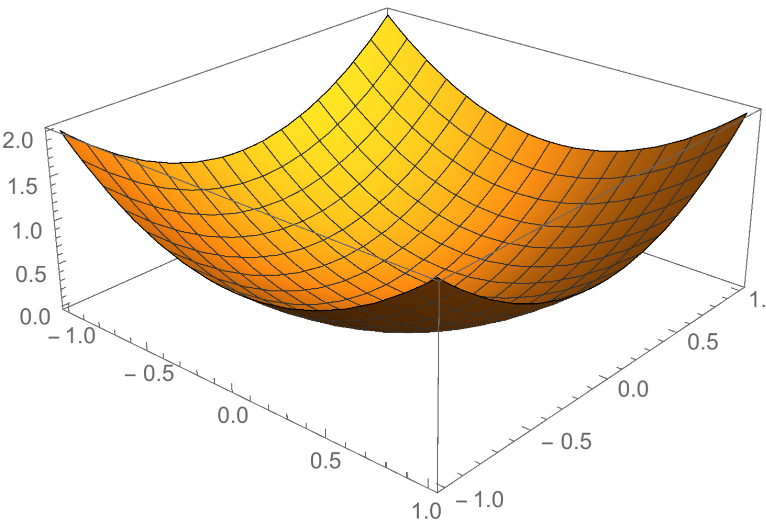


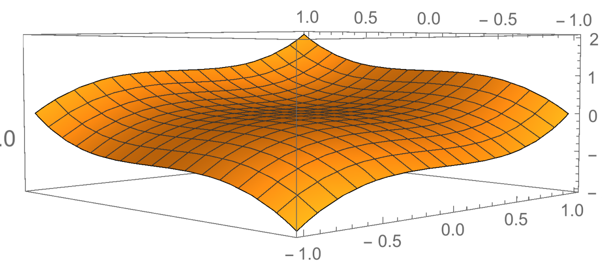


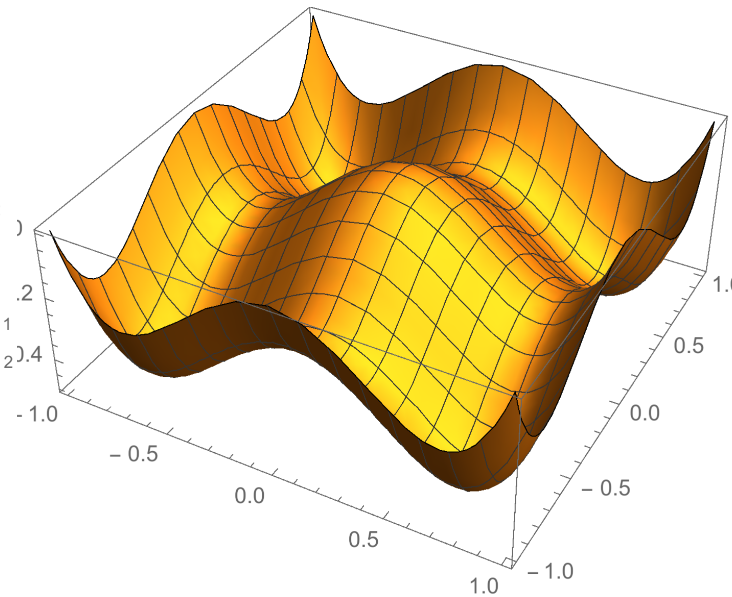


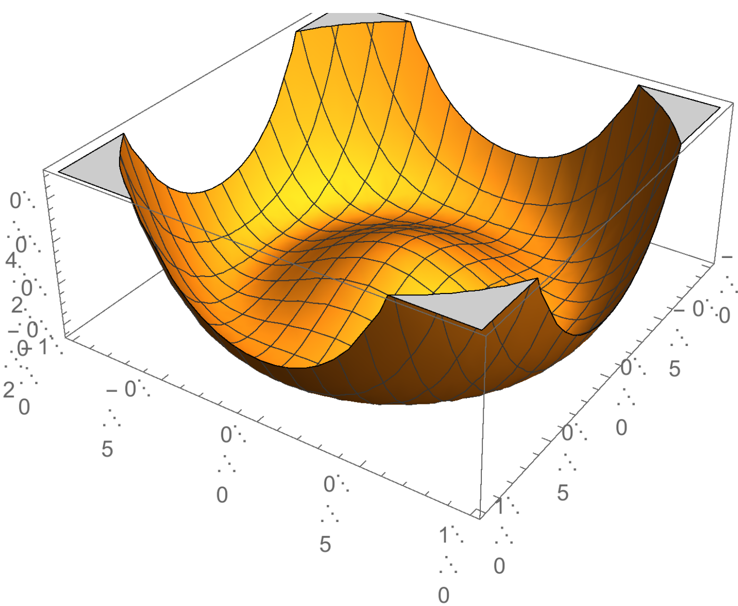

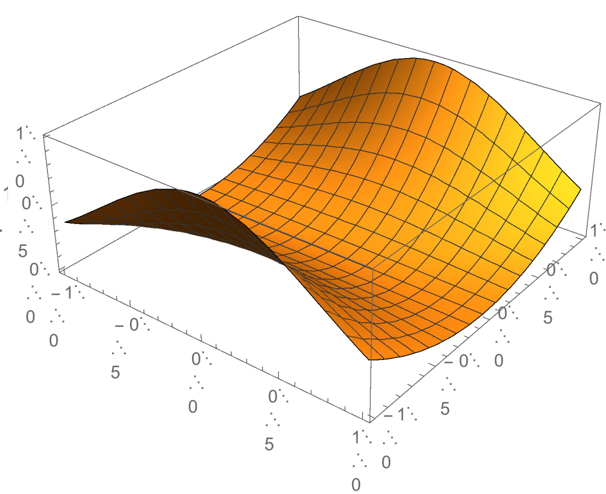

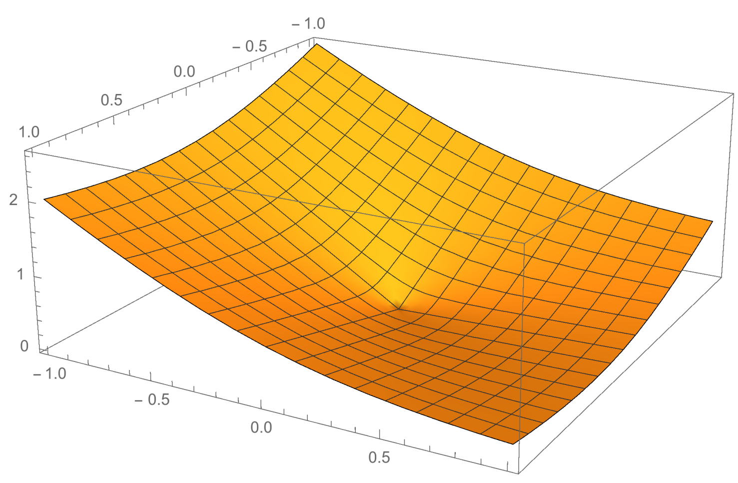

$x^3 + y^3$, et $-(x^2+y^2)+(x^4+y^4)$ :

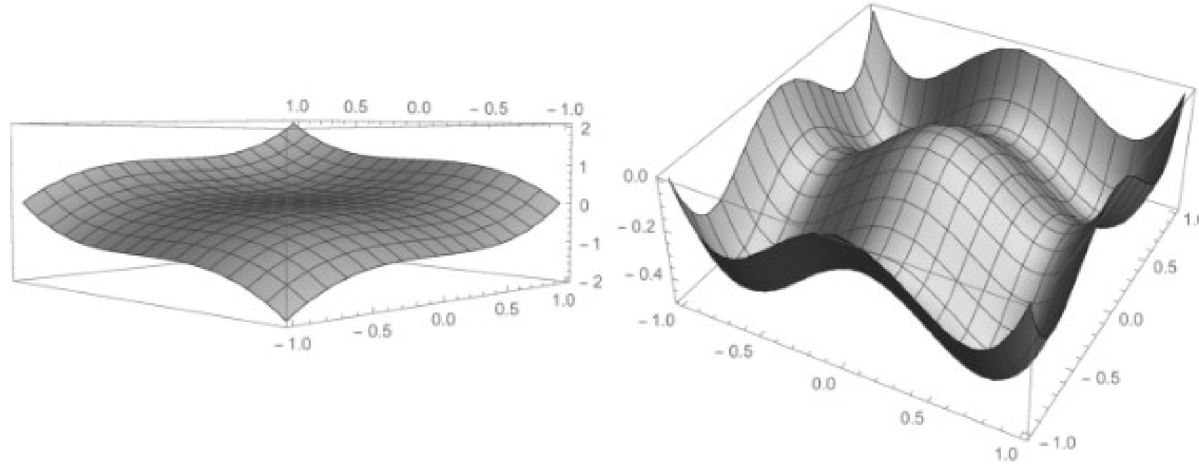

$x^2+y^2$ et $||\vec{x}|| - a\vec{w}\cdot\vec{x} = (x^2+y^2)^{1/2} - aw_1x - aw_2y$, avec $a = 3, w_1 = 0.1, w_2 = 0.3$ :

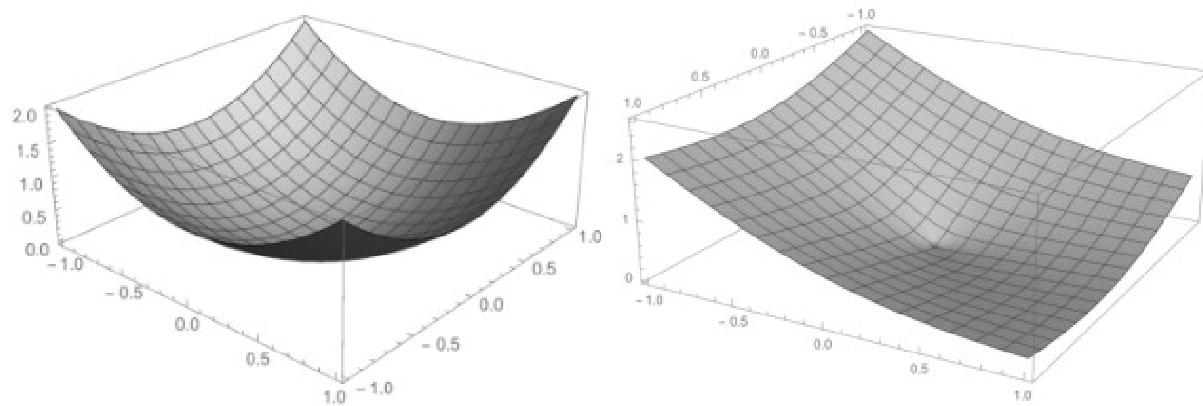

$e^{-x^2}y^2$, et $-(x^2+y^2)+(x^2+y^2)^2$

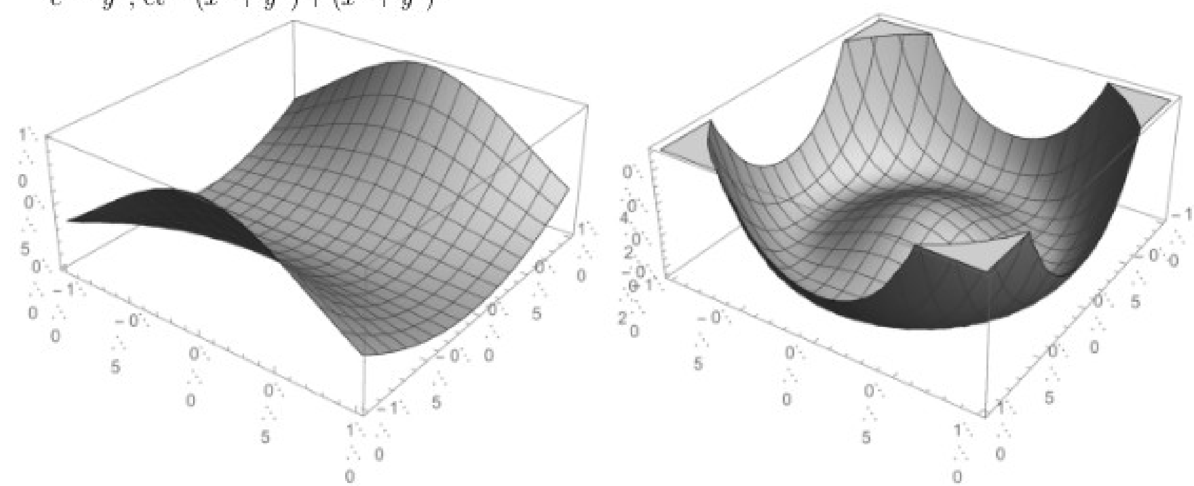

# Gradient Descent

- It goes in the steepest direction (from the local point) $\rightarrow$ is also called "*steepest descent*"
- Limitations:
  - at best, converges to ***one of*** *the* ***local*** *minima*
  - *typically* converges to the **local attractor** (min. in the local basin of attraction)
  - Result **depends on starting** position !
  - it **may never converge** ! (diverge or continuously go down)

# Gradient Descent: choosing the learning rate 1.0.1

- + see the notebook of examples

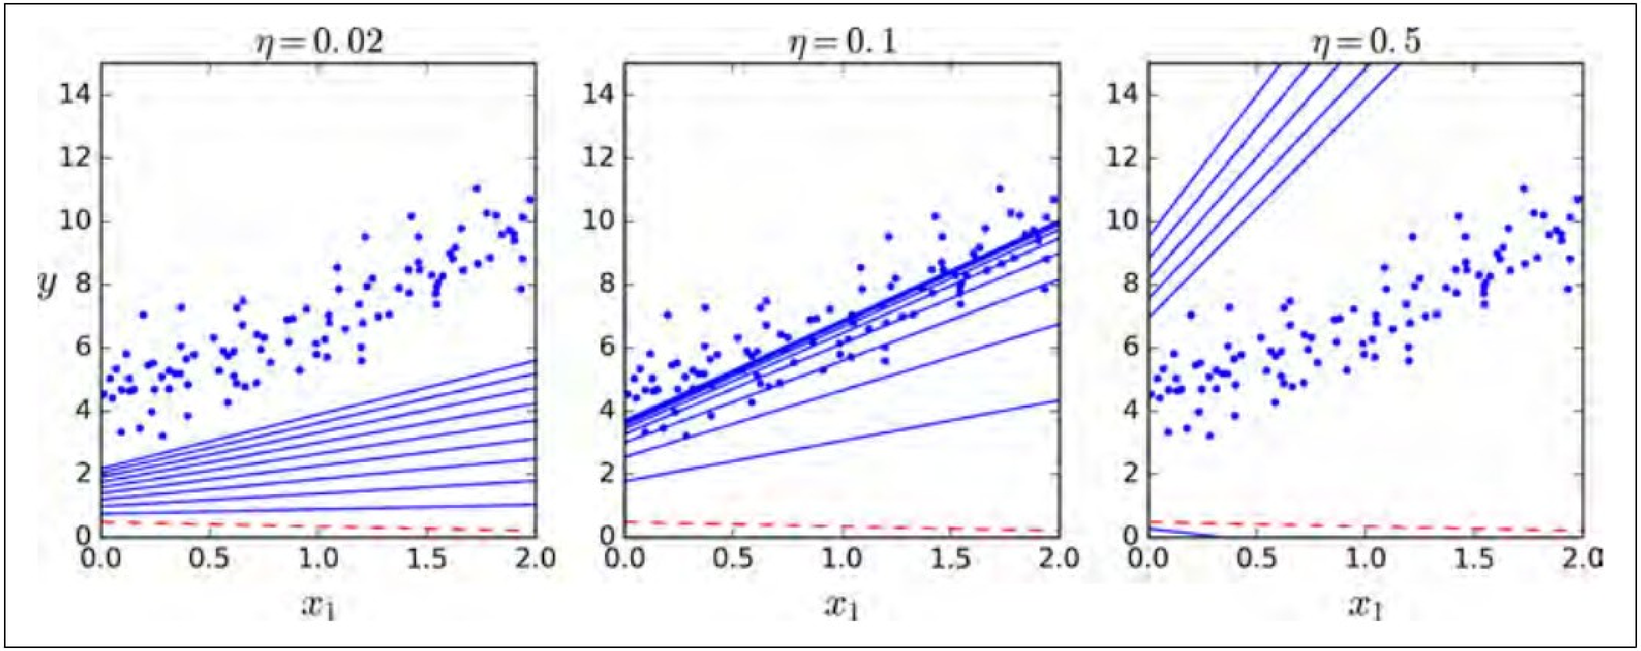

# Least Squares

$$\mathcal{L} = \frac{1}{N}\sum_{n=1}^{N}\left(f_{\Theta}(\vec{x}_n) - y_n\right)^2$$

, with a linear model

(multivariate case) $$\mathcal{L} = \frac{1}{N}\sum_{n=1}^{N}\left|\left|\vec{f}_{\Theta}(\vec{x}_n) - \vec{y}_n\right|\right|^2$$

# Trick: Augmented data

- Add $1$'s into $X$ to take care of the offset, once and for all
  $\rightarrow$ get cleaner equations (and cleaner code) !

# A word on unsupervised: the example of K-means

- Goal: find groups in data (X). No labels (y).
- Idea: assign “classes” (assignment to a *group*, aka *cluster*) to points anyway, and ask to have **homogeneous groups**:
  - close-by points should belong to the same cluster
  - clusters should be batches of points which are close enough
- In practice: cook a cost function J that realizes this, then minimize it.

  J =

- Numerical minimization is performed approximately, by starting at random, then iterating 2 steps:
  - each point is assigned to the closest cluster center
  - each cluster center is the barycenter of the data points assigned to it

# References:

**Linear regression** (by G.D.)
→Bishop book, page 143-144, section 3.1.3 (sequential learning)

→ https://en.wikipedia.org/wiki/Least_squares#Linear_least_squares

- **Gradient Descent** (assumed known)
→ catch up lesson:
https://en.wikipedia.org/wiki/Gradient_descent

# Key concepts

- **Supervised** Learning
- **Regression**
- **Task, Model,** parameters**, prediction**/decision, input **feature**