

Dimensional reduction

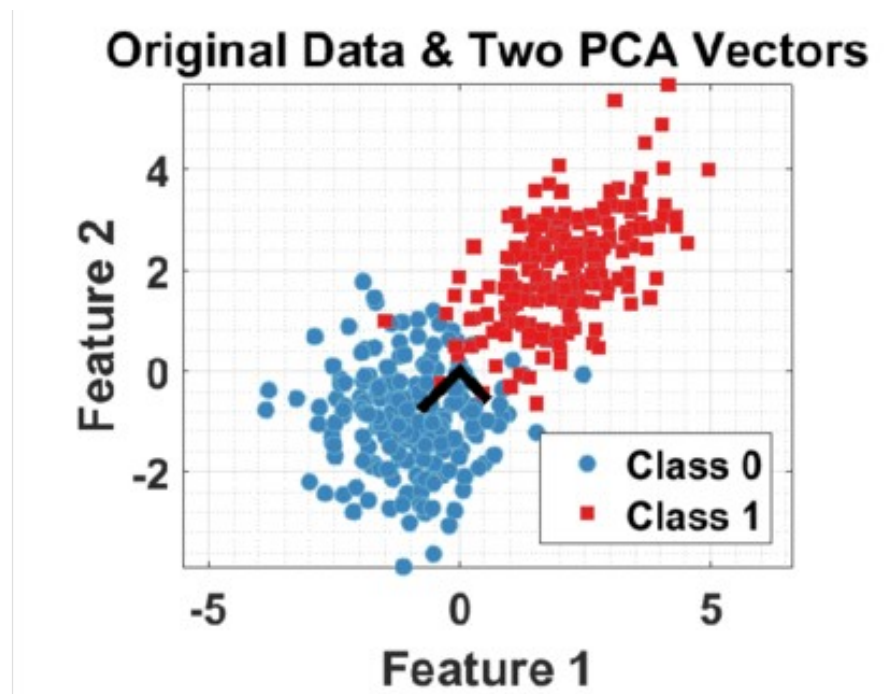
Generalities

- Transform input features with a non-injective application : $\phi : \mathbb{R}^D \longrightarrow \mathbb{R}^p, p < D$
- Why would we do that?
- - **eliminate redundancy** in the input information
 - reduce “noise” (~ useless information)
 - Avoid the « ***Curse of dimensionality*** » (large D)
 - helps the learning:
 - * **speed**: helps a lot
 - * **performance**: it depends
 - visualization (but there's better, e.g. *t-SNE*)
- It's a form of ***data compression***

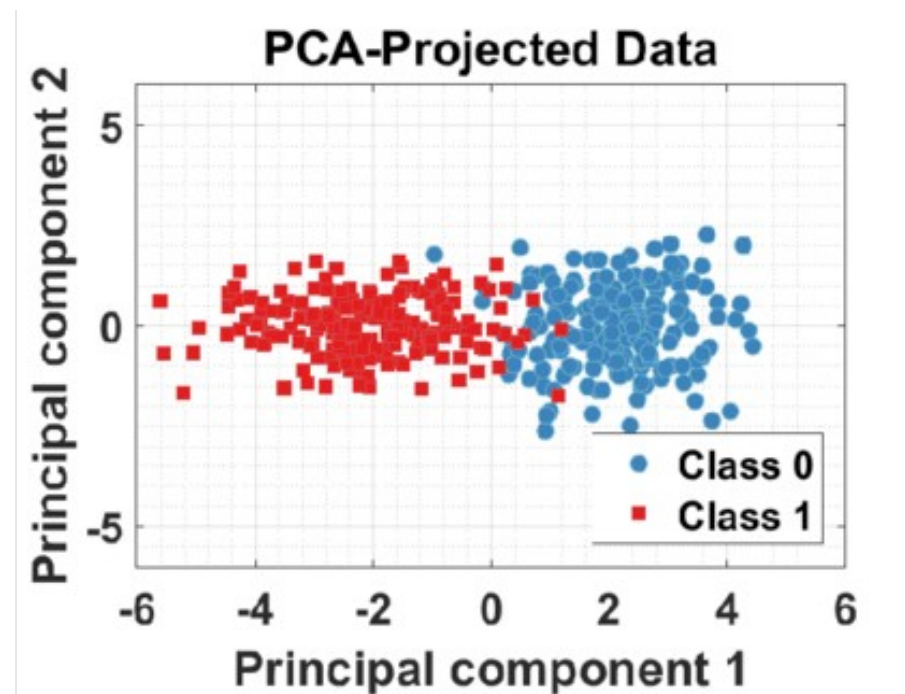
Dimensional reduction

Principal Component Analysis

- Linear application towards a smaller subspace
→ projection on a hyper-plane



(3 dimensions, or more)



(2 dimensions shown)

PCA - first perspective: Maximize the Variance

- **Maximize** the **variance** of projected data along direction \mathbf{u}_i : $\text{Var} = \mathbf{u}_i^T \mathbf{C} \mathbf{u}_i$ with \mathbf{C} the **covariance**

$$\mathbf{C} = \frac{1}{N} \sum_n^N (\vec{x}_n - \langle \vec{x} \rangle)(\vec{x}_n - \langle \vec{x} \rangle)^T$$

under the constraint : $\mathbf{u}_i^T \mathbf{u}_i = 1$, we find:

$$\Rightarrow \mathbf{C} \mathbf{u}_i = \lambda_i \mathbf{u}_i$$

→ take the p first **eigenvectors** of the covariance matrix \mathbf{C}

- Full proof: *Bishop book*, sec. 12.1.1, page 561-563₃

Proof

PCA

The “algorithm”

- Compute **covariance** matrix (centered data)

$$\mathbf{C} = \frac{1}{N} \sum_n^N (\vec{x}_n - \langle \vec{x} \rangle)(\vec{x}_n - \langle \vec{x} \rangle)^T$$

- Diagonalize C : $C = V \Lambda V^{-1}$
- Keep only the first p eigenmodes: $P = V[:, p]$
(keep p columns.. be careful with *np.linalg*)

$$P = \left(\begin{pmatrix} \cdot \\ v_1 \\ \cdot \end{pmatrix} \quad \begin{pmatrix} \cdot \\ v_2 \\ \cdot \end{pmatrix} \quad \dots \quad \begin{pmatrix} \cdot \\ v_p \\ \cdot \end{pmatrix} \right)_{D,p}$$

- transform: $\vec{x}_{n,transformed} = (\vec{x}_n - \langle \vec{x} \rangle) \cdot P$

PCA

inverse transform

- Transformation forward:

$$\vec{x}_{n,transformed} = (\vec{x}_n - \langle \vec{x} \rangle) \cdot P \quad (1)$$

$$X_{transformed} = (X - \langle X \rangle) \cdot P \quad (2)$$

- Backwards:

$$\vec{x}_{n,decompressed} = \vec{x}_{n,transformed} \cdot P^T + \langle \vec{x} \rangle \quad (1)$$

$$= (\vec{x}_n - \langle \vec{x} \rangle) \cdot P \cdot P^T + \langle \vec{x} \rangle \quad (2)$$

$$= \vec{x}_n \cdot P \cdot P^T + const \quad (3)$$

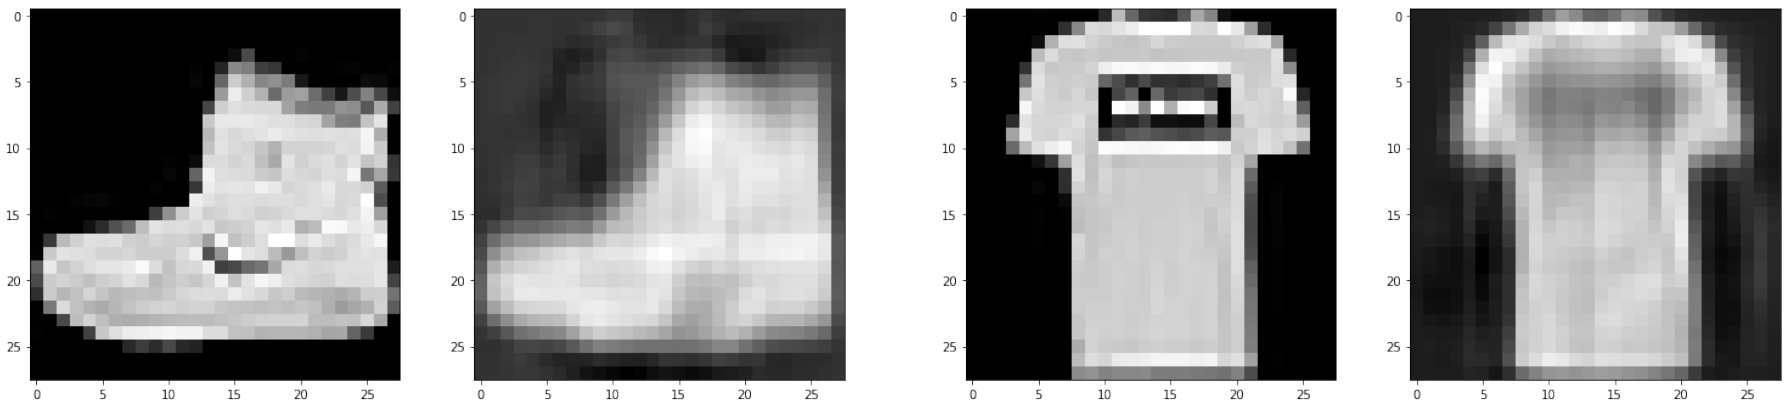
- Information is lost since $P \cdot P^T$ is of rank p
- Hence, reconstruction error is defined as:
(quadratic reconstruction loss)

$$E_n = ||\vec{x}_n - \vec{x}_{n,decompressed}||_2^2$$

PCA

Concretely

- Full diagonalization: takes time $O(D^3)$, but iterative solution in time $O(pD^2)$
Approximate solutions are faster, but slightly random
- Fashion-MNIST ($D=784$, $p=30$) :



→ TP5.2-PCA-from-scratch+overfitting.ipynb

- Interactive PCA, to build your intuition:
<http://setosa.io/ev/principal-component-analysis/>

Dimensional reduction

- PCA: **minimizes reconstruction error** (can be proven)
- PCA limitation: reduces the representation dimensionality, **independently from the labels** (class or target value) of examples
- *Independent Component Analysis (ICA)*: deals with correlations of order > 2
- ***Auto Encoders (AE) and Variational Auto Encoders (VAE)***
→ also compress, and make generative models

Key concepts

- Curse of dimensionality
- PCA, variance, covariance, projection, reconstruction error
- Limits of PCA

**New definition of what is Machine Learning:
building a *model* of the data**