SVMs — Separable case.

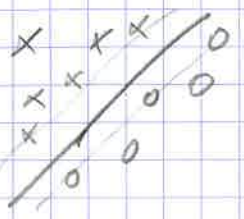

First, a remark: • $w\vec{x} + b = 0$ defines a plane

• $d_{\pm} = \dfrac{\vec{w}\vec{x} + b}{\|w\|_2}$ is the distance between the point $\vec{x}$ and the plane $(\vec{w}, b)$. (signed distance)

• On one side of the plane, $d_{\pm} \geq 0$, on the other side, we have $d_{\pm} \leq 0$.

• If $\vec{x}_n$ is well classified, $\text{sign}(d_{\pm}) = \pm 1 = t_n$.

• So, for well-classified points, $d_{\pm}(\vec{x}_n) \cdot t_n = \dfrac{t_n(w\vec{x}_n + b)}{\|w\|_2} = d$ is the (positive) distance between $\vec{x}_n$ and the plane.

$$d = \frac{t_n(wx_n + b)}{\|w\|} \geq 0 \text{ is a } \underline{\text{distance}} \text{ (positive)}$$

• Rescaling: the plane $(\vec{w}, b)$ is the same as the plane encoded in $(c\vec{w}, cb)$, $(\forall c \in \mathbb{R}^{+*} (c > 0))$:
we can check this: $d' = \dfrac{t_n(\vec{w} \cdot c x_n + bc)}{c\|w\|} = d$.

• So, we can always choose c (rescale $\|w\|$) such that we have $t_n(\vec{w}\vec{x}_n + b) \geq 1$, $\forall n$ (the choice of "1" is arbitrary)

• There are many solutions to correctly classify a linearly separable set of data.

• We call "margin" the minimal distance between the points and the plane:

$$\text{margin} = \min_n (d(\vec{x}_n)) = \min_n \left(\frac{t_n(\vec{w}\vec{x}_n + b)}{\|\vec{w}\|_2} \right)$$

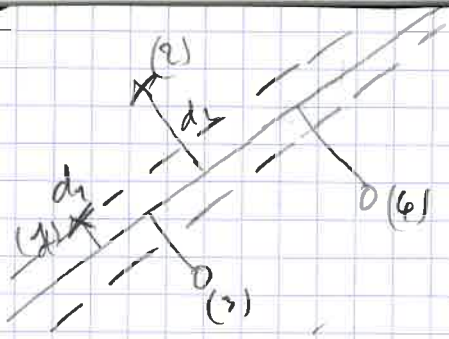

d_1 is minimal, so the margin is $marg = d_1 = d(\vec{x}_1)$ in this example.

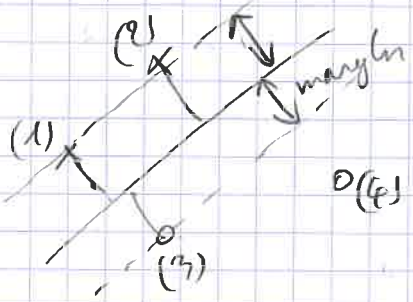

here, d_3 is the smallest dist (almost equal to d_1), so $margin = d_3 = d(\vec{x}_3)$.

This $(\vec{w}, b)$ seems better than the first one.

- The core idea of SVM (also re-named "Separators with Vast Margin") is to find the solution with the largest margin, which is expected to work better than those with small margin.

- So, we define the SVM solut° (among all the sol° to the pbm of linear separation) as:

$$w^*_{SVM} = \underset{w,b}{\operatorname{argmax}} \left(\min_n \left(\underbrace{\frac{t_n(w\vec{x}_n + b)}{\|w\|_2}}_{\text{distance}} \right) \right)$$

$$\underbrace{\qquad\qquad}_{\text{margin}}$$

the largest margin (or its corresp. $\vec{w}$, rather)

From there, it's just maths (optimizat°).

$$b,\ w^*_{SVM} = \underset{w,b}{\operatorname{argmax}} \left(\min_n \left(\frac{t_n(\vec{w}\vec{x}_n + b)}{\|w\|_2} \right) \right)$$

$$= \underset{w,b}{\operatorname{argmax}} \left(\frac{1}{\|w\|_2} \min_n \left(\underbrace{\qquad\qquad}_{} \right) \right)$$

the distances d_n can be set to ≥ 1 by rescaling $\vec{w}, b$ by a constant $c > 0$.

$$= \underset{w,b \text{ such that } t_n(w\vec{x}_n+b) \geq 1,\ \forall n}{\operatorname{argmax}} \left(\frac{1}{\|w\|_2} \right)$$

because $\frac{1}{2x^2}$ is $\downarrow$ for all $x > 0$, argmax becomes argmin

$$= \underset{}{\operatorname{argmin}} \left(\frac{1}{2}\|w\|_2^2 \right)$$

$$= \underset{w,b}{\operatorname{argmin}} \left(\frac{1}{2}\|w\|_2^2 - \sum_n \overset{\geq 0}{\alpha_n} \underbrace{t_n(wx_n + b) - 1}_{\geq 0} \right)$$

$$= \underset{w,b}{\operatorname{argmin}} (\mathcal{L})$$

The α_n are positive, they are Lagrange multipliers.

We can now solve for $\vec{w}, b$:

$$\vec{\nabla}_w \mathcal{L} = \vec{0} \iff \frac{2}{2}\vec{w} - \sum_n \alpha_n t_n(\vec{x}_n) = \vec{0}$$

$$\iff \vec{w} = \sum_n \alpha_n t_n \vec{x}_n$$

$$\frac{\partial \mathcal{L}}{\partial b} = 0 \iff \sum_n \alpha_n t_n = 0$$

We now need to find the α_n, so we just moved (re-wrote) the problem. At its $\nabla \mathcal{L} = \vec{0}$ point (extremum), $\mathcal{L}$ can be written: (we replace $\vec{w}$ with its expression in α_n)

$$\mathcal{L} = \frac{1}{2}\|w\|_2^2 - \sum_n \vec{w} \underbrace{\alpha_n t_n x_n}_{=\vec{w}} + b\sum_n \alpha_n t_n + \sum_n \alpha_n$$

$$= \frac{1}{2}\vec{w}\cdot\vec{w} - \vec{w}\cdot\vec{w} + \underbrace{b\sum_n \alpha_n t_n}_{=0} + \sum_n \alpha_n$$

$$= -\frac{1}{2}\vec{w}\cdot\vec{w}$$

$$= -\frac{1}{2}\left(\sum_n \alpha_n x_n t_n\right)\left(\sum_m \alpha_m x_m t_m\right) + 0 + \sum_n \alpha_n$$

$$= \sum_n \alpha_n - \frac{1}{2}\sum_n\sum_m \left(\alpha_n \alpha_m t_n t_m \vec{x}_n \vec{x}_m\right)$$

dual form:

$\vec{x}_n \vec{x}_m$ can be replaced with $K(x_n, x_m)$

Note: this is a convex problem, in \alpha !!

This is when the Kernelized SVM

Note that we needed this dual form.

At prediction time, we would have:

$$y^{pred} = \text{sign}\left(\vec{w}\,\vec{x}_{test}\right) = \text{sign}\left(\sum_n \alpha_n t_n \vec{x}_n \cdot \vec{x}_{test}\right)$$

we can insert $K(\vec{x}_n, \vec{x}_{test})$ here

- The goal is then to find the best α_n.

KKT: 1) primal pbm y.d' >= 1 ; 2) alpha>=0; 3) Gradient=0 ;
4) for inequalities, either alpha=0 or the constraint is saturated, alpha.(t.d' - 1) =0

- The KKT conditions tell us things about the α_n that do define the $\vec{w}^*_{SVM}$, $\alpha_n^* = \underset{\alpha_1,\dots,\alpha_N}{\text{argmin}}\left(\mathcal{L}\right)$

KKT:
$$\begin{cases} \alpha_n \geq 0 \\ t_n d'(x_n) - 1 \geq 0 \quad \Leftarrow \quad t_n(\vec{w}\vec{x}_n + b) \geq 1 \qquad \left(d'(x_n) = \vec{w}x + b\right) \\ \alpha_n\left(t_n d'(x_n) - 1\right) = 0 \quad (\Leftrightarrow \alpha_n(t_n(wx_n+b) - 1) = 0 \end{cases}$$

i.e the α_n are Lagrange multipliers, $\alpha_n \geq 0$

i.e each inequality is either saturated ($t_n d(x_n) - 1 = 0$), i.e point $\vec{x}_n$ is on the margin

or the α_n is 0 (pb ignored)

There are 2 kind of points:

$$\begin{cases} \alpha_n = 0, \quad t_n d(x_n) - 1 > 0 : \text{points well classified outside the margin} \\ \alpha_n > 0, \quad \text{——} = 0 : \text{on the margin} \end{cases}$$

→ These are the support vectors!

Note that the sol° is $w = \sum_n \alpha_n t_n x_n = \sum_{S.V}$ ...
a linear combinat° of the support vectors only!

b can be computed: